

From Prediction to Support: An Equity-Centered Adaptive Learning Analytics Framework for Workforce Readiness

Abstract

Community college workforce programs are under pressure to help learners persist through credentials while preparing for changing skill demands. Learning analytics and adaptive learning systems can support this work, but prediction alone does not create equitable or actionable support. This article develops the Equity-Centered Adaptive Pathways (ECAP) Framework for responsible adaptive learning analytics in community college workforce programs. Rather than reporting a field deployment or model-validation study, the article uses an illustrative pre-implementation scenario to show how institutional teams can reason about learner signals, support thresholds, intervention intensity and staff capacity before piloting analytics-driven supports. The framework connects learner-state sensing, pedagogical interpretation, adaptive support pathways, workforce-readiness alignment, human-centered governance and resource calibration. It translates current learning analytics, adaptive learning, guided pathways, responsible learning analytics and educational AI governance literature into practical guidance for instructional designers, advisors, analytics teams and community college leaders. The article argues that analytics are most useful when connected to transparent, equitable and feasible support pathways.

Keywords: adaptive learning analytics, community colleges, educational technology, learning analytics, pre-implementation planning, responsible learning analytics, student success, workforce readiness, workforce upskilling

1 Introduction

Community colleges are central to U.S. workforce development because they connect access, short-cycle credentials, transfer pathways, noncredit education and local employer needs. The American Association of Community Colleges reported that community colleges served 10.6 million learners from fall 2024 to fall 2025, including 6.5 million credit and 4.1 million noncredit students ([American Association of Community Colleges, 2026](#)). Recent enrollment patterns reinforce that role. The National Student Clearinghouse Research Center reported that community college certificate enrollments reached 752,000 in fall 2025, a 28.3% increase from fall 2021 ([National](#)

[Student Clearinghouse Research Center, 2026](#)). Workforce pressure is also increasing. The U.S. Bureau of Labor Statistics projects 5.2 million additional jobs from 2024 to 2034, with growth driven mainly by healthcare and social assistance ([U.S. Bureau of Labor Statistics, 2025](#)). Globally, the World Economic Forum projects substantial skill disruption by 2030 and reports that skill gaps are a major barrier to organizational transformation ([World Economic Forum, 2025](#)).

These numbers matter because community college workforce learners often move through programs while working, caregiving, commuting, using uneven technology access, changing schedules and making high-stakes decisions about credentials and employment. Faculty, advisors and program leaders are asked to improve completion, support adult learners, align programs with employer needs and use data more effectively. Guided pathways reforms have shown the value of clearer program maps, integrated advising and whole-college support structures, but those reforms also require careful coordination across people, systems and milestones ([Bailey, Jaggars, & Jenkins, 2015](#); [Minaya & Acevedo, 2024](#)). Yet many institutions already have more data than they can interpret well. Learning management systems, early-alert tools, advising platforms, assessment systems and dashboards can identify patterns, but they do not automatically create educationally meaningful action.

This creates a practical educational technology problem for practitioners: how can institutions move from learner data to timely support without reducing students to risk scores, overburdening staff or deepening inequities through opaque automation? Learning analytics has long been described as a way to understand and improve learning environments, but scholars continue to warn that analytics must remain about learning rather than only analytics ([Gašević, Dawson, & Siemens, 2015](#); [Wise, Knight, & Ochoa, 2021](#)). Dashboard reviews show increasing attention to self-regulated learning, feedback and recommendations, but also uneven maturity and limited evidence of sustained implementation ([Matcha, Uzir, Gašević, & Pardo, 2020](#); [Paulsen & Lindsay, 2024](#)). Adaptive learning reviews likewise suggest promise, while emphasizing that effects depend on design, context, pedagogy and implementation quality ([du Plooy, Casteleijn, & Franzsen, 2024](#); [Kulik & Fletcher, 2016](#)).

The risk is that predictive analytics can look precise while remaining educationally thin. A missed assignment, low login frequency or weak quiz score may indicate confusion, work conflicts, caregiving demands, broadband barriers, assessment anxiety or a course-design artifact. Each interpretation calls for a different support response. If the system simply labels a learner as high risk, it may make the institution faster at noticing a problem without making it better at helping the student.

This article develops the Equity-Centered Adaptive Pathways (ECAP) Framework as a response to that gap. The framework treats adaptive learning analytics as a support-decision pathway rather than an automation pipeline. It integrates learner-state sensing, pedagogical interpretation, adaptive support pathways, workforce-readiness alignment, human-centered governance and resource calibration. Its novelty lies less in proposing another predictive model and more in specifying how analytics can be designed as a transparent, equitable and feasible pathway for community college workforce programs.

The article is a framework and practice article with an illustrative scenario analysis. It does not report a field deployment, causal evaluation or validated machine-learning

model. Instead, it uses a bounded pre-implementation scenario to show how institutional teams can reason about support thresholds, intervention intensity, workload and scenario boundaries before real learners are affected. This genre choice is deliberate. Community college leaders often need a defensible way to prepare pilots before they have enough local data for a formal impact study. Instructional designers need to know what support resources must exist before recommendations are turned on. Advisors and faculty need to know when an alert is a helpful prompt and when it is noise. Students need transparency, choice and ways to correct the record. The contribution is therefore threefold: a research-grounded ECAP framework, concrete implementation artifacts for practitioners and a scenario-planning approach that makes assumptions and capacity tradeoffs visible before pilot implementation.

2 Theoretical and Practical Foundations

The ECAP framework draws from five bodies of literature that are useful for community college workforce programs: self-regulated learning, situated and competency-oriented learning, learning analytics and dashboards, adaptive learning, and responsible sociotechnical implementation.

Self-regulated learning emphasizes how learners plan, monitor, seek help, use feedback and adjust strategies (Panadero, 2017; Pintrich, 2004; Zimmerman, 2002). This matters for adaptive analytics because many student-facing dashboards and recommendation systems are intended to help learners notice progress and make decisions. Reviews of learning analytics dashboards show that self-regulated learning is a common theoretical foundation, but they also show that dashboards often remain more descriptive than actionable (Matcha et al., 2020; Paulsen & Lindsay, 2024). ECAP therefore treats learner signals as prompts for feedback, reflection and support rather than as labels that define the learner. The distinction is especially important for adult and part-time learners. A dashboard that merely displays lagging indicators may confirm that a learner is falling behind after the opportunity for timely support has passed. A support-oriented pathway, by contrast, asks what feedback, planning resource, schedule adjustment or help-seeking option could still change the learner’s next step.

Situated learning and communities of practice shift attention from course activity alone to participation in meaningful practice (Brown, Collins, & Duguid, 1989; Lave & Wenger, 1991). Workforce readiness is not only completion of assignments; it involves competencies, credential momentum, tool use, safety practices, communication, professional judgment and alignment with program goals. Guided pathways reforms make a similar institutional argument: students benefit when programs are structured, milestones are clear, advising is integrated and support is connected to program maps (Bailey et al., 2015; Jenkins, Lahr, & Brock, 2024; Jenkins, Lahr, & Mazzariello, 2021). Recent community college research also emphasizes that guided pathways become meaningful only when reforms are implemented across advising, instruction, equity work and student-success routines rather than as isolated technical tools (Minaya & Acevedo, 2024). For community college workforce programs, analytics should therefore connect learner support to pathways and competencies, not only to generic retention indicators.

Learning analytics and adaptive learning research provide the technological foundation, but they also caution against technology-first claims. Learning analytics can support sensemaking, feedback and intervention when it is integrated into educational practice (Ferguson, 2012; Knight, Gibson, & Shibani, 2020; Siemens, 2013). The field has increasingly asked not only whether analytics can predict, but whether learning analytics research matters for learning, design and decision making in authentic settings (Blumenstein, 2020; Wise et al., 2021). Adaptive learning can personalize resources and pacing, but its effects are mediated by instructional design, learner context and implementation conditions (du Plooy et al., 2024; Kulik & Fletcher, 2016; Luo, Zheng, Yin, & Teo, 2025; VanLehn, 2011). For this reason, ECAP treats prediction as one step in a larger support workflow. The framework does not assume that a more complex model is automatically a better educational tool. It asks whether the signal can be interpreted, whether the response is pedagogically appropriate and whether support can be delivered without creating new burdens.

Responsible learning analytics literature adds the ethical foundation. Scholars have emphasized privacy, consent, transparency, student agency, vulnerability and institutional power relationships (Jones, 2019; Pardo & Siemens, 2014; Prinsloo & Slade, 2016; Slade & Prinsloo, 2013). Ethical guidance in learning analytics has also moved toward practical checklists and design processes that help teams discuss data sources, legitimacy, access and intervention before deployment (Drachsler & Greller, 2016; Prieto-Alvarez, Martinez-Maldonado, & Shum, 2020). Algorithmic bias research shows that educational data systems can reproduce inequities if historical patterns are treated as neutral evidence (Baker & Hawn, 2022), and recent fairness work in multimodal learning analytics reinforces that fairness, accountability, transparency and ethics must be handled across data collection, modeling and intervention design (Jin et al., 2024). Human-centered learning analytics and AI research reinforces the need for stakeholder involvement, human control, safety, trustworthiness and contextual design (Alfredo et al., 2024; Holstein, McLaren, & Aleven, 2019; Topali et al., 2025). ECAP therefore makes governance and student agency core design requirements.

Finally, implementation science and educational design research explain why a technically plausible system can fail in practice. Implementation outcomes such as acceptability, feasibility, fidelity, cost and sustainability are distinct from learner outcomes and should be considered before scale-up (Proctor et al., 2011). Educational design research similarly treats theory and practice as mutually informing and emphasizes iterative development in real contexts (McKenney & Reeves, 2018). These traditions support the use of pre-implementation scenarios as planning tools: institutions can surface assumptions, estimate workload and test governance logic before launching analytics-driven supports.

3 Equity-Centered Adaptive Pathways Framework

The Equity-Centered Adaptive Pathways Framework begins from a simple premise: adaptive learning analytics becomes educationally meaningful only when it helps people make better support decisions. The framework connects six components into a support-decision chain. Learner-state sensing notices possible need. Pedagogical interpretation




Fig. 1 Equity-Centered Adaptive Pathways Framework. The framework connects learner-state sensing with pedagogical interpretation, adaptive support, workforce-readiness alignment, human-centered governance and resource calibration. The outer guardrails emphasize transparency, fairness, privacy and student agency.

asks what the signal may mean in context. Adaptive support pathways connect the interpretation to feasible action. Workforce-readiness alignment keeps the action tied to credentials, competencies and career goals. Human-centered governance keeps the system transparent, contestable and fair. Resource calibration keeps the support model feasible for the institution.

Figure 1 presents ECAP as a pathway rather than a pipeline. This distinction is important. A pipeline implies that data flow toward an automated decision. A pathway implies that signals are interpreted, reviewed and connected to support. If any component is removed, adaptive analytics loses educational meaning, workforce relevance, equity safeguards or implementation feasibility.

The pathway also changes who should be involved in design. A predictive pipeline can be designed mainly by technical teams and then handed to practitioners. ECAP requires shared design from the beginning because each component depends on local knowledge. Faculty know which milestones actually reflect learning. Advisors know how students explain registration, attendance and schedule disruptions. Instructional designers know whether a recommended resource is accessible and aligned with course goals. Institutional research teams know what data are reliable enough to use. Students know whether outreach feels respectful or stigmatizing. These perspectives are central to the evidence needed to make analytics usable.

Table 1 turns the framework into a practitioner review tool. The mapping draws on self-regulated learning, situated learning, guided pathways, adaptive learning, responsible learning analytics and implementation research (Brown et al., 1989; du Plooy et al., 2024; Jenkins et al., 2021; Lave & Wenger, 1991; Proctor et al., 2011; Slade & Prinsloo,

Table 1 ECAP Design Questions for Practitioner Review

ECAP component	Practitioner question	Design move	Risk if neglected
Learner-state sensing	Which signals suggest a timely support need without labeling the learner?	Use milestones, assessment progress, participation and help-seeking as prompts for review.	Surveillance, proxy scoring or deficit labeling.
Pedagogical interpretation	What might this signal mean in this course and learner context?	Discuss plausible explanations before selecting support.	Generic or punitive outreach that misses the real need.
Adaptive support pathways	What support follows, and who owns the next action?	Route learners to practice, tutoring, advising, feedback, peer support or schedule review.	Alerts identify need but do not create action.
Workforce-readiness alignment	How does support connect to credentials, competencies or career goals?	Link progress signals to program maps, competency evidence and credential milestones.	Generic retention metrics are mistaken for workforce readiness.
Human-centered governance	Who can understand, review, contest and revise recommendations?	Use plain-language explanations, human review, appeal paths and fairness audits.	Black-box sorting, privacy harm, distrust or inequitable escalation.
Resource calibration	Can the institution deliver the support volume and intensity implied?	Estimate alert volume, workload, escalation timing and advising or tutoring capacity.	Alert fatigue, delayed support and unsustainable workload.

2013; Zimmerman, 2002). Each component exists because it addresses a recognizable failure mode in analytics-enabled support.

3.1 Design Principles for Practice

Five design principles translate ECAP into practical decisions. First, institutions should prioritize interpretability before optimization. A slightly less complex signal that educators can understand may be more useful in an early pilot than a stronger black-box model. Second, support should come before prediction. Every alert should be paired with a feasible support action, an owner and a student-facing explanation.

Third, workforce relevance should be precise. A course-progress indicator should not be described as employability. Institutions should distinguish credential momentum, competency evidence, licensure preparation, work-based learning and employment outcomes. Fourth, equity should be built into the design rather than appended as a review step. Data minimization, fairness auditing, consent or notice, appeal routes and student agency belong in the workflow from the start (Jones, 2019; Prinsloo &

[Slade, 2016](#); [UNESCO, 2023](#)). Fifth, adaptation must be capacity-aware. Thresholds and support intensity should be chosen with faculty, advisors, tutoring centers and students, because an alert that no one can act on is not support.

These principles make ECAP more than a diagram. They define the conditions under which analytics can become a responsible educational technology practice: the system must be interpretable, actionable, workforce-relevant, equitable and feasible.

For practitioners, the framework can be used as a design review tool. A team can take any proposed alert and ask six questions: What signal triggered it? What are the plausible explanations? What support action follows? How does the action relate to a credential or competency milestone? What transparency, privacy and appeal procedures protect students? What workload will the support create? If the team cannot answer these questions, the alert is not ready for a pilot even if the underlying model appears technically strong.

4 Framework Development and Illustrative Scenario Design

This article uses framework development with an illustrative pre-implementation scenario. The purpose is not to validate an intervention or compare predictive models. The purpose is to show how an institutional team could use ECAP to reason about support logic, workload and governance before piloting analytics-driven supports with real students.

4.1 Framework Development

The framework was developed by synthesizing literature on learning analytics, adaptive learning, self-regulated learning, situated learning, guided pathways, responsible learning analytics, educational AI governance and implementation. The synthesis asked four practice-oriented questions. What learner signals might indicate support needs? How should educators interpret those signals? What support actions are feasible in community college workforce programs? What governance and capacity conditions must be in place before scale-up?

The resulting ECAP components were then organized as a support-decision pathway. This approach follows educational design research by using theory to shape practical design artifacts, then using design artifacts to make the theory more actionable ([McKenney & Reeves, 2018](#)). It also follows implementation science by treating feasibility, acceptability, workload and sustainability as early design concerns rather than post-hoc explanations for failure ([Proctor et al., 2011](#)).

4.2 Evidence Sources

The manuscript uses three evidence sources. Public reports establish the community college and workforce context. Peer-reviewed literature grounds the framework and practical recommendations. Archived local CSV artifacts provide scenario inputs for a bounded planning exercise. [Table 2](#) summarizes how each source type is used.

Table 2 Evidence Sources and Use in the Pre-Implementation Design

Source type	Use in manuscript	Interpretation limits
Official higher education and workforce reports	Establish community college scale, certificate growth, workforce demand and skills-change context	Used for context only; not merged with learner-level records
Peer-reviewed learning analytics, adaptive learning and ethics literature	Grounds ECAP in self-regulation, sociotechnical analytics, adaptive support, implementation and responsible data use	Supports design principles and practice guidance, not local effect sizes
Archived scenario inputs	Provide illustrative assumptions for support-threshold and intervention-intensity planning	Planning artifacts only; not raw training data, field outcomes or validated model outputs

4.3 Illustrative Scenario Boundaries

The scenario was intentionally placed after the framework because its role is illustrative. It uses archived scenario inputs to ask what an institutional team might need to discuss before a pilot: What happens if support thresholds are too low or too high? How much staff capacity is required as support intensity increases? Which outcome proxies are meaningful, and which require stronger local definition? Table 3 makes these boundaries explicit.

The scenario does not use identifiable student records, raw learner-level data, deployed intervention data, saved model objects or reproducible prediction files. The values function as planning artifacts, not observed outcomes. This boundary reduces the risk of treating a design scenario as empirical proof and keeps the article focused as a framework and practice contribution.

The scenario is still useful because many implementation failures begin before a pilot launches. Teams may choose a threshold without asking how many alerts it will create. They may define workforce readiness with a convenient course proxy rather than with faculty-validated competencies. They may assume that advisors can absorb extra outreach without reducing other services. They may also assume that students will understand a recommendation without seeing the data logic behind it. A bounded scenario allows these assumptions to be discussed while the stakes are still low. In that sense, scenario analysis is not a substitute for evidence; it is a preparation for collecting better evidence responsibly.

4.4 Interpretation Boundaries and Governance Assumptions

Scenario estimates are reported as percentage points when they compare percentages. They are not reported with inferential tests, confidence intervals or causal language because the archived artifacts do not support those claims. The scenario assumes that every analytics-triggered support would be reviewed by a human, communicated in plain language and monitored for fairness and workload. Any live pilot would require institutional review, privacy and compliance analysis, student notice or consent procedures, data-retention rules, appeal processes and continuous monitoring for bias

Table 3 Scenario Assumptions and Boundaries

Scenario choice	Planning purpose	Not evidence of
Static pathway	Represents a common baseline with uniform pacing and reactive support.	Observed control-group outcomes.
Adaptive pathway	Tests how interpreted learner signals might trigger reviewed support.	A deployed intervention.
Threshold setting	Makes alert volume, timing and staff review needs visible.	An optimized policy rule.
Support intensity	Helps teams estimate workload, escalation and coordination needs.	Proof that more intensive support is always better.
Workforce-readiness proxy	Prompts local definition of workforce-relevant milestones.	Job placement, wages or employer evaluation.

and drift ([Jones, 2019](#); [National Institute of Standards and Technology, 2023](#); [Pardo & Siemens, 2014](#)).

5 Illustrative Scenario Analysis for Pilot Planning

The scenario analysis illustrates how ECAP can support pilot planning. The important point is not the exact percentage difference in any scenario. The important point is that adaptive learning analytics requires explicit assumptions about who is flagged, why they are flagged, what support follows, who reviews the action and whether the institution can deliver the support responsibly.

5.1 Pathway Assumptions

The archived scenario compared static-support and adaptive-support assumptions. Under the modeled assumptions, the adaptive-support condition generated higher planning estimates for completion, retention and workforce-readiness alignment than the static-support condition. These estimates are planning artifacts for local discussion rather than evidence of improved student outcomes.

The planning value is that the comparison forces a team to define what “adaptive support” actually means. If an adaptive pathway is only a different dashboard label, the estimate has little practical value. If the pathway includes earlier feedback, advising outreach, skill practice, tutoring referral, schedule review and human review, then the institution can begin to estimate what resources and governance are required.

This is where ECAP differs from a conventional model-comparison narrative. The relevant design question is not which algorithm produces the most favorable number in an archive. The relevant question is whether a learner signal can be converted into support that a student can understand and an institution can deliver. For example, a low score in an early technical module might call for practice resources if the issue is skill fluency, an instructor conversation if the issue is misunderstanding, an advising check-in if the issue is schedule pressure or no alert at all if the score reflects a known




Fig. 2 Illustrative support-threshold and intervention-intensity scenario. Values are archived planning inputs used to reason about capacity calibration, not observed learner outcomes.

assessment-design problem. The same numeric signal can lead to different actions depending on context.

5.2 Threshold and Intensity Tradeoffs

Support thresholds shape who receives attention and when. A lower threshold may identify learners earlier, but it can also produce more alerts than staff can address. A higher threshold may reduce alert volume, but it can miss learners who would benefit from early support. The threshold question is therefore not only technical. It is also pedagogical, ethical and operational.

Intervention intensity creates a similar tradeoff. More frequent or personalized support may be desirable, but it requires advising capacity, faculty time, tutoring availability, accessible learning resources and communication systems. The scenario assigns progressively larger planning estimates as intensity increases, but those values should be read as a prompt for capacity discussion, not as evidence that higher intensity will be better in every setting.

Figure 2 is therefore best read as a design conversation starter. It asks institutional teams to decide how many alerts can be reviewed each week, which supports are available, when escalation is appropriate and what stop rules should prevent overreach. It also encourages equity review: if one group is consistently overflagged or receives less benefit from support, the institution should pause and revise the pathway before scale-up.

A practical reading of the figure is to ask what would happen on Monday morning after the system is turned on. If the chosen threshold creates more alerts than advisors can review, students may receive delayed or generic outreach. If the pathway calls for tutoring but tutoring has limited hours, the recommendation may frustrate students rather than help them. If increased intensity means repeated messages without student choice, the intervention may feel like surveillance. Conversely, if thresholds are too restrictive, support may arrive only after a learner has already withdrawn or missed a

credential milestone. The figure therefore supports workload and governance planning, not claims about optimal policy.

5.3 What the Scenario Does and Does Not Support

The scenario supports pre-implementation reasoning about assumptions, workload and governance. It does not support claims about predictive accuracy, causal impact, field effectiveness or workforce outcomes. Conventional model-performance metrics are outside the article’s main claim and are not reported as findings in the main text. The framework would still require field testing with students, faculty, advisors, instructional designers and program leaders.

This evidence posture is important for the article’s framework-and-practice contribution. The contribution is not that a model achieved a performance threshold. The contribution is a practical framework and a set of design artifacts that can help community college teams prepare safer, more transparent pilots before analytics-driven supports affect real learners.

6 Practical Guidance for Community College Workforce Programs

ECAP is intended to be used by cross-functional teams, not only analytics specialists. A useful pilot team should include program faculty, advisors, instructional designers, student-success staff, institutional research or analytics staff, privacy or compliance representatives, program leaders and student representatives when feasible. If formal student representation is not feasible, teams should still test messages, support choices and data-use explanations with students before launch. Each group sees a different part of the pathway.

The starting point should be a real program problem rather than a data feature. A team might choose a first-term health care certificate, an advanced manufacturing sequence or an information technology bootcamp because the program has clear milestones, known support resources and faculty willing to review the pathway. Beginning with one program keeps the pilot small enough to learn from. It also makes the workforce-readiness construct more concrete because competencies, schedules, lab requirements and credential gates differ across programs.

6.1 From Risk Scores to Support Decisions

The first practical step is to create an indicator-to-support matrix before configuring alerts. Table 4 shows how a team can connect signals to possible interpretations, support actions, responsible roles, student-facing messages, escalation rules, equity checks and capacity notes. This artifact helps prevent alert fatigue because every signal must earn its place in the workflow.

6.2 Role-Specific Responsibilities

Instructional designers should translate analytic signals into learning resources tied to course milestones and accessibility needs. Advisors should treat alerts as prompts

Table 4 Indicator-to-Support Matrix for Adaptive Pathway Pilots

Learner signal		Likely interpretation	Support response	Equity and capacity check
Missed milestone	early	Confusion, work conflict, access issue or unclear directions.	Advisor sends a supportive nudge with a recovery option and milestone guide.	Check overflagging patterns; limit alerts to high-value milestones.
Low performance	assessment	Skill gap, test anxiety, weak prior knowledge or misaligned practice.	Instructor offers targeted practice, worked examples, tutoring or feedback.	Compare support uptake and benefit across learner groups.
Reduced activity	LMS	Work schedule, technology access, disengagement or course-design artifact.	Success coach sends a nonjudgmental check-in and verifies access.	Avoid alerts where LMS activity is not meaningful.
Credential-plan mismatch		Progress no longer aligns with program map or career goal.	Advisor reviews schedule, prerequisites and credential options.	Review effects on part-time students and align with advising cycles.
Weak competency evidence	compe-	Task completion is not producing workforce-relevant skill evidence.	Faculty adds skill practice, lab support or an employer-aligned task.	Use only faculty-validated competencies and check assessment bias.
Repeated support non-response	nonre-	Message timing, trust concern, competing obligations or wrong channel.	Staff try another channel and invite student-selected support.	Use human review, avoid penalty and set weekly contact caps.

for conversation rather than labels, document the support offered and escalate only when need or repeated nonresponse warrants it. Analytics teams should prioritize interpretable indicators, monitor subgroup patterns and report drift or overflagging before scale-up. Program leaders should assign ownership, protect staff capacity, define stop rules and keep workforce-readiness claims tied to locally defined competencies. Students should know what data are used, why support is being offered, what choices they have and how to correct or contextualize the record.

Student-facing communication should be plain and nondeficit. A useful pattern is: “We noticed [specific progress signal]. This can happen for several reasons. Here are [two support options]. You can also tell us if the signal does not reflect your situation.” Messages such as “you are high risk” or “you are unlikely to complete” should be avoided because they turn an analytic signal into an identity label.

The same principle applies to staff-facing language. Advisors and instructors should not receive a black-box score with no explanation. They need a short rationale, the signal that triggered the recommendation, the support options already available and any known limitations. For example, an alert might say that a learner missed a required safety-practice checkpoint and has not accessed the make-up module, not that the

Table 5 Pilot Readiness Checklist for Responsible Adaptive Learning Analytics

Readiness area	Question	Evidence needed	Stop or proceed indicator
Data readiness	Are indicators accurate, current and meaningful?	Data dictionary, missing-data review and faculty validation.	Proceed only if indicators are interpretable and auditable.
Privacy readiness	What data are used, who sees them and how long are they retained?	Privacy review, access rules, retention plan and student-facing notice.	Stop if data use cannot be clearly explained to students.
Human review workflow	Who reviews alerts before action?	Workflow map, reviewer roles and override process.	Proceed only if review is assigned and documented.
Advisor capacity	How many alerts can staff address each week?	Workload estimate and response-time target.	Pause if alert volume exceeds response capacity.
Faculty buy-in	Do indicators reflect meaningful milestones?	Faculty review notes and course milestone map.	Stop if alerts reflect artifacts rather than learner needs.
Student communication	Do students receive respectful explanations and choices?	Message templates and student usability feedback.	Proceed only after students understand the purpose.
Fairness audit plan	Who may be overflagged, underflagged or underserved?	Subgroup monitoring plan and benefit-distribution metrics.	Pause if subgroup patterns suggest inequitable treatment.
Escalation and appeal	How can students correct or contextualize data?	Appeal process and correction workflow.	Stop if no correction or appeal route exists.
Workload monitoring	Does the pilot add sustainable work?	Weekly workload logs and staff feedback.	Pause if staff cannot complete support actions on time.
Stop or pause criteria	What requires redesign before scale-up?	Predefined review thresholds.	Scale only after support is useful, fair and feasible.

learner is unlikely to finish the program. This wording keeps attention on the next support decision.

6.3 Pilot Readiness

Before scale-up, institutions should test whether the pathway is feasible. Table 5 provides a pilot readiness checklist aligned with implementation outcomes such as feasibility, fidelity and sustainability (Proctor et al., 2011). A team should be willing to pause if any readiness area lacks an owner, evidence base or stop rule.

6.4 Using the Framework in One Program

A practical ECAP pilot can begin in one workforce program with clear milestones. The team should define three to five interpretable signals, map each signal to a support option, estimate weekly alert volume, draft student messages, identify who reviews alerts and schedule a midpoint fairness and workload review. The pilot should measure

process outcomes before claiming learner outcomes: whether alerts were understandable, whether support was delivered, whether students perceived the outreach as useful and whether staff could sustain the workflow.

A simple pilot timeline can keep the work manageable. During the first month, the team defines the program milestones, reviews data quality and drafts the indicator-to-support matrix. During the second month, faculty, advisors and students review message templates and support options. During the third month, the team runs a shadow pilot in which alerts are generated but not used for automated outreach; this allows staff to inspect false alarms, missing cases and workload. Only after that review should the team begin limited live outreach. At the end of the term, the team should decide whether to continue, narrow, redesign or stop the pilot based on feasibility, student experience, fairness patterns and support delivery.

7 Equity, Governance and Student Agency

Adaptive learning analytics is most defensible when it is governed as an equity intervention. Without safeguards, predictive systems can reproduce historical inequities, increase surveillance or make learners feel judged by a hidden system (Baker & Hawn, 2022; Prinsloo & Slade, 2016; Slade & Prinsloo, 2013). ECAP therefore treats governance as part of the framework rather than as an administrative appendix.

7.1 Governance Responsibilities

Responsible governance begins with role clarity. Faculty and advisors interpret learner signals in context. Instructional designers ensure that support resources are usable, accessible and aligned with learning goals. Analytics teams document indicators, monitor quality and audit subgroup patterns. Program leaders decide whether staffing and resources can support the intervention. Compliance and privacy staff review data use, access rules and retention. Students should be represented in the design and should receive explanations that make support transparent.

Governance should also specify what the system is not allowed to do. ECAP should not silently restrict course options, steer students away from credentials, trigger punitive action or share sensitive inferences beyond the support team. It should not use demographic variables to assign individual risk. Demographic and equity-related variables may be needed for fairness auditing, but that is different from using them to decide which student receives which opportunity. The governance team should document these boundaries before launch so that the pilot has a clear ethical perimeter.

Table 6 summarizes safeguards that should be operational before deployment. These safeguards align with broader responsible AI guidance emphasizing transparency, privacy, explainability, fairness, accountability and human oversight (National Institute of Standards and Technology, 2023; UNESCO, 2023; U.S. Department of Education, Office of Educational Technology, 2023). In educational settings, these values must be translated into everyday procedures: who sees the alert, who contacts the student, what the student is told, how inaccurate data are corrected and when the system is paused. ECAP should also avoid inferred sensitive traits, continuous monitoring beyond

Table 6 Responsible Implementation Safeguards for Adaptive Learning Analytics

Risk area	Safeguard	Practical implementation
Opaque prediction	Explainability	Alerts include the educational reasons for support, not only a risk score
Privacy and surveillance	Data minimization	Use only data needed for support decisions and publish student-facing data-use notices
Bias and equity drift	Fairness auditing	Monitor flagging, support delivery and benefit across relevant learner groups
Over-automation	Human oversight	Instructors or advisors review alerts and can override recommendations
Student agency	Transparent communication	Students receive support options, explanations and ways to contextualize records
Workload overload	Resource calibration	Thresholds and support intensity are matched to staffing and reviewed each term

educationally necessary signals and secondary use of data beyond support delivery unless those uses are separately reviewed and disclosed.

7.2 Fairness and Workload Monitoring

Fairness audits should examine both flagging and benefit. A group may be overflagged because the data reflect structural barriers. A group may be under-supported because outreach channels do not work equally well. A group may receive support but not benefit if the support is poorly matched to need. ECAP therefore requires monitoring of who is flagged, who receives support, how quickly support is delivered and whether support is useful. Subgroup analysis should use locally appropriate categories, protect small groups from re-identification and examine intersectional patterns where sample size allows.

Workload monitoring is also an equity issue. If staff cannot respond to alerts promptly, students with fewer outside resources may be most harmed by delayed support. Institutions should track weekly alert volume, response time, unresolved cases and staff feedback. If the pathway creates hidden burden for advisors, faculty or student-success staff, the responsible decision may be to narrow the pilot rather than expand it.

Fairness monitoring should be interpreted cautiously. A subgroup difference in flagging rates is not automatically evidence of unfairness, because different learners may face different course schedules, work demands or program milestones. But subgroup differences should always prompt investigation. The team should ask whether the underlying indicator is educationally meaningful, whether missing data are unevenly distributed, whether the support offered is equally accessible and whether the observed benefit is similar across groups. This turns fairness auditing into a design conversation rather than a compliance formality.

7.3 Student Agency

Student agency requires more than notice. Students should understand why support is offered, what data informed the outreach, what options they have and how they can correct or contextualize a signal. They should not be silently moved into lower-opportunity pathways because of an analytic prediction. They should be invited into support conversations that preserve choice and dignity. In workforce programs, this is especially important because learners may be balancing education with employment, caregiving and urgent financial constraints that a data system cannot fully see.

The governance goal is not to make students trust a model. The goal is to make institutional support worthy of trust.

Student agency also requires attention to refusal and preference. A student may decline a particular form of support, prefer a different communication channel or explain that a data signal reflects a temporary circumstance. The system should allow that information to change the support pathway. Disputed or corrected records should not trigger punitive escalation, lower-opportunity routing or lasting labels while review is pending. Otherwise, analytics can become rigid even when the language around it sounds supportive. A human-centered pathway should remember that students are not data points moving through a workflow; they are people making decisions under constraints that institutions may only partially observe.

8 Limitations and Future Design-Based Research

This article has clear boundaries. The illustrative scenario is synthetic and pre-implementation. It is not field validation, does not establish causal impact and does not represent all community college contexts. The scenario values are useful for planning conversations about thresholds, intensity and capacity, but they should not be interpreted as observed improvements in completion, retention or workforce readiness.

The workforce-readiness construct also requires local definition. In some programs, readiness may involve licensure preparation, lab competencies, safety procedures or clinical performance. In others, it may involve tool mastery, communication, portfolio evidence, internship participation or confidence in next steps. ECAP requires institutions to define workforce-readiness indicators with program faculty, advisors, students and employer-informed evidence rather than relying on generic proxies.

The framework itself has not yet been tested with community college students, advisors, faculty, instructional designers or program leaders. Future research should proceed in layers. First, participatory design sessions should test whether the six ECAP components reflect actual support decisions. Second, data-governance and fairness reviews should determine which indicators are appropriate and which should be excluded. Third, prototype testing should examine whether messages and recommendations are understandable, respectful and useful. Fourth, a small pilot should use human review, documented support actions and workload monitoring. Fifth, mixed-method evaluation should examine process outcomes, learner outcomes, staff experience and equity patterns.

The strongest future evidence would combine quantitative and qualitative data. Quantitative data can show whether learners who receive support persist, complete

or demonstrate competencies at different rates. Qualitative data can show whether learners experienced support as timely and respectful. Staff interviews can reveal whether the pathway improved coordination or simply added another task. These forms of evidence would help determine whether ECAP functions as a human-centered support system rather than a prediction layer attached to existing practice.

Design-based research is especially appropriate because ECAP is a framework for situated practice rather than a fixed product. A future study could begin with co-design workshops, produce a local pathway prototype, run a shadow-pilot analysis, revise the intervention and then examine a small live pilot. Each iteration would test both theory and usability. For example, researchers could ask whether learner-state sensing improves only after faculty redefine milestones, whether students understand the data-use notice, whether advisors can act within the expected response window and whether support produces similar benefits across learner groups. These questions are more useful at this stage than asking whether a generic model is universally accurate.

Finally, institutions should resist premature scale-up. A pilot should expand only after the support pathway is useful, fair and feasible. This means the institution can explain the data use, deliver support on time, monitor subgroup patterns, correct inaccurate records and pause the system when harms or overload appear.

9 Conclusion

Community college workforce programs need educational technologies that help learners persist toward meaningful credentials and competencies. Prediction alone is insufficient for that work. A risk score may identify a pattern, but it does not explain the learner’s situation, select an appropriate support, protect student agency or ensure that staff can respond.

The Equity-Centered Adaptive Pathways Framework contributes a practical way to connect analytics with support. It links learner-state sensing, pedagogical interpretation, adaptive support pathways, workforce-readiness alignment, human-centered governance and resource calibration. The framework asks institutions to treat analytics as a support pathway that must be understandable, equitable and feasible.

The illustrative scenario in this article shows how teams can reason about thresholds, intensity and capacity before a live pilot. Its value is not empirical proof. Its value is that it makes assumptions visible and helps teams ask better implementation questions before students are affected. Field validation remains necessary, but pre-implementation reasoning can make future pilots safer and more useful.

For instructional designers, advisors, analytics teams and community college leaders, the central message is direct: adaptive learning analytics should improve support decisions. Its promise lies not in prediction itself, but in helping institutions notice needs earlier, interpret them carefully and respond with transparent supports tied to learner goals and workforce pathways.

References

Alfredo, R., Echeverria, V., Jin, Y., Yan, L., Swiecki, Z., Gašević, D., Martinez-Maldonado, R. (2024). Human-centred learning analytics and AI in education:

- A systematic literature review. *Computers and Education: Artificial Intelligence*, 6, 100215, <https://doi.org/10.1016/j.caeai.2024.100215>
- American Association of Community Colleges (2026). *AACC Fast Facts 2026*. https://www.aacc.nche.edu/wp-content/uploads/securepdfs/2026/03/AACC2026_Fast_Fact_Sheet.pdf. (Accessed 17 May 2026)
- Bailey, T., Jaggars, S.S., Jenkins, D. (2015). *Redesigning america's community colleges: A clearer path to student success*. Harvard University Press.
- Baker, R.S., & Hawn, A. (2022). Algorithmic bias in education. *International Journal of Artificial Intelligence in Education*, 32, 1052–1092, <https://doi.org/10.1007/s40593-021-00285-9>
- Blumenstein, M. (2020). Synergies of learning analytics and learning design: A systematic review of student outcomes. *Journal of Learning Analytics*, 7(3), 13–32, <https://doi.org/10.18608/jla.2020.73.3>
- Brown, J.S., Collins, A., Duguid, P. (1989). Situated cognition and the culture of learning. *Educational Researcher*, 18(1), 32–42, <https://doi.org/10.3102/0013189X018001032>
- Drachsler, H., & Greller, W. (2016). Privacy and analytics: It's a DELICATE issue: A checklist for trusted learning analytics. *Proceedings of the sixth international conference on learning analytics & knowledge* (pp. 89–98).
- du Plooy, E., Casteleijn, D., Franzsen, D. (2024). Personalized adaptive learning in higher education: A scoping review of key characteristics and impact on academic performance and engagement. *Heliyon*, 10(21), e39630, <https://doi.org/10.1016/j.heliyon.2024.e39630>
- Ferguson, R. (2012). Learning analytics: Drivers, developments and challenges. *International Journal of Technology Enhanced Learning*, 4(5/6), 304–317, <https://doi.org/10.1504/IJTEL.2012.051816>
- Gašević, D., Dawson, S., Siemens, G. (2015). Let's not forget: Learning analytics are about learning. *TechTrends*, 59(1), 64–71, <https://doi.org/10.1007/s11528-014-0822-x>

- Holstein, K., McLaren, B.M., Aleven, V. (2019). Co-designing a real-time classroom orchestration tool to support teacher–AI complementarity. *Journal of Learning Analytics*, 6(2), 27–52, <https://doi.org/10.18608/jla.2019.62.3>
- Jenkins, D., Lahr, H., Brock, T. (2024). *Lessons from two major evaluations of guided pathways reforms* (Tech. Rep.). Community College Research Center. Retrieved from <https://files.eric.ed.gov/fulltext/ED646216.pdf>
- Jenkins, D., Lahr, H., Mazzariello, A. (2021). *How to achieve more equitable community college student outcomes: Lessons from six years of CCRC research on guided pathways* (Tech. Rep.). Community College Research Center. Retrieved from <https://ccrc.tc.columbia.edu/publications/equitable-community-college-student-outcomes-guided-pathways.html>
- Jin, Y., Echeverria, V., Yan, L., Zhao, L., Alfredo, R., Tsai, Y.S., ... Martinez-Maldonado, R. (2024). FATE in MMLA: A student-centred exploration of fairness, accountability, transparency, and ethics in multimodal learning analytics. *Journal of Learning Analytics*, 11(3), 6–23, <https://doi.org/10.18608/jla.2024.8351>
- Jones, K.M.L. (2019). Learning analytics and higher education: A proposed model for establishing informed consent mechanisms to promote student privacy and autonomy. *International Journal of Educational Technology in Higher Education*, 16, 24, <https://doi.org/10.1186/s41239-019-0155-0>
- Knight, S., Gibson, A., Shibani, A. (2020). Implementing learning analytics for learning impact: Taking tools to task. *The Internet and Higher Education*, 45, 100729, <https://doi.org/10.1016/j.iheduc.2020.100729>
- Kulik, J.A., & Fletcher, J.D. (2016). Effectiveness of intelligent tutoring systems: A meta-analytic review. *Review of Educational Research*, 86(1), 42–78, <https://doi.org/10.3102/0034654315581420>
- Lave, J., & Wenger, E. (1991). *Situated learning: Legitimate peripheral participation*. Cambridge University Press.
- Luo, J., Zheng, C., Yin, J., Teo, H.H. (2025). Design and assessment of AI-based learning tools in higher education: A systematic review. *International Journal of Educational Technology in Higher Education*, 22, 42, <https://doi.org/10.1186/s41239-025-00540-2>

- Matcha, W., Uzir, N.A., Gašević, D., Pardo, A. (2020). A systematic review of empirical studies on learning analytics dashboards: A self-regulated learning perspective. *IEEE Transactions on Learning Technologies*, 13(2), 226–245, <https://doi.org/10.1109/TLT.2019.2916802>
- McKenney, S., & Reeves, T.C. (2018). *Conducting educational design research* (2nd ed.). Routledge.
- Minaya, V., & Acevedo, N. (2024). *Whole-college reforms in community colleges: Guided pathways practices and early academic success in three states* (Tech. Rep. No. CCRC Working Paper No. 136). Community College Research Center. Retrieved from <https://eric.ed.gov/?id=ED646255>
- National Institute of Standards and Technology (2023). *Artificial intelligence risk management framework (AI RMF 1.0)* (Tech. Rep.). National Institute of Standards and Technology.
- National Student Clearinghouse Research Center (2026). *Clearinghouse final fall enrollment trends*. <https://nscresearchcenter.org/final-fall-enrollment-trends/>. (Accessed 17 May 2026)
- Panadero, E. (2017). A review of self-regulated learning: Six models and four directions for research. *Frontiers in Psychology*, 8, 422, <https://doi.org/10.3389/fpsyg.2017.00422>
- Pardo, A., & Siemens, G. (2014). Ethical and privacy principles for learning analytics. *British Journal of Educational Technology*, 45(3), 438–450, <https://doi.org/10.1111/bjet.12152>
- Paulsen, L., & Lindsay, E. (2024). Learning analytics dashboards are increasingly becoming about learning and not just analytics: A systematic review. *Education and Information Technologies*, 29(11), 14279–14308, <https://doi.org/10.1007/s10639-023-12401-4>
- Pintrich, P.R. (2004). A conceptual framework for assessing motivation and self-regulated learning in college students. *Educational Psychology Review*, 16, 385–407, <https://doi.org/10.1007/s10648-004-0006-x>
- Prieto-Alvarez, C.G., Martinez-Maldonado, R., Shum, S.B. (2020). LA-DECK: A card-based learning analytics co-design tool. *Lak20 conference proceedings: The tenth international conference on learning analytics & knowledge* (pp. 63–72).

- Prinsloo, P., & Slade, S. (2016). Student vulnerability, agency and learning analytics: An exploration. *Journal of Learning Analytics*, 3(1), 159–182, <https://doi.org/10.18608/jla.2016.31.10>
- Proctor, E., Silmere, H., Raghavan, R., Hovmand, P., Aarons, G., Bunger, A., ... Hensley, M. (2011). Outcomes for implementation research: Conceptual distinctions, measurement challenges and research agenda. *Administration and Policy in Mental Health and Mental Health Services Research*, 38, 65–76, <https://doi.org/10.1007/s10488-010-0319-7>
- Siemens, G. (2013). Learning analytics: The emergence of a discipline. *American Behavioral Scientist*, 57(10), 1380–1400, <https://doi.org/10.1177/0002764213498851>
- Slade, S., & Prinsloo, P. (2013). Learning analytics: Ethical issues and dilemmas. *American Behavioral Scientist*, 57(10), 1510–1529, <https://doi.org/10.1177/0002764213479366>
- Topali, P., Ortega-Arranz, A., Rodríguez-Triana, M.J., Er, E., Khalil, M., Akçapınar, G. (2025). Designing human-centered learning analytics and artificial intelligence in education solutions: A systematic literature review. *Behaviour & Information Technology*, 44(5), 1071–1098, <https://doi.org/10.1080/0144929X.2024.2345295>
- UNESCO (2023). *Guidance for generative AI in education and research* (Tech. Rep.). United Nations Educational, Scientific and Cultural Organization. Retrieved from <https://www.unesco.org/en/articles/guidance-generative-ai-education-and-research>
- U.S. Bureau of Labor Statistics (2025). *Employment projections: 2024–34*. <https://www.bls.gov/emp/>. (Accessed 17 May 2026)
- U.S. Department of Education, Office of Educational Technology (2023). *Artificial intelligence and the future of teaching and learning: Insights and recommendations* (Tech. Rep.). U.S. Department of Education. Retrieved from <https://www2.ed.gov/documents/ai-report/ai-report.pdf>
- VanLehn, K. (2011). The relative effectiveness of human tutoring, intelligent tutoring systems and other tutoring systems. *Educational Psychologist*, 46(4), 197–221, <https://doi.org/10.1080/00461520.2011.611369>

Wise, A.F., Knight, S., Ochoa, X. (2021). What makes learning analytics research matter. *Journal of Learning Analytics*, 8(3), 1–9, <https://doi.org/10.18608/jla.2021.7647>

World Economic Forum (2025). *The future of jobs report 2025*. <https://www.weforum.org/publications/the-future-of-jobs-report-2025/>. (Accessed 17 May 2026)

Zimmerman, B.J. (2002). Becoming a self-regulated learner: An overview. *Theory Into Practice*, 41(2), 64–70, https://doi.org/10.1207/s15430421tip4102_2